

IN3200/IN4200: Dissecting Home Exam 1


Spring 2021

Objectives of home exam 1

- Hands-on experience with two common data storage schemes
- Implementation (and testing) of serial functions
- OpenMP parallelization
- Reflections on performance related issues

Connectivity graph

- Nodes represent individual data objects
- Edges represent *direct-connections* between pairs of “nearest neighbors”
- Symmetry: if node u is a nearest neighbor of node v , then v is a nearest neighbor of u



2D table as storage format for connectivity graph

A 2D table of values 0 and 1 (implemented as a 2D array of type `char**`)

| | | | | | |
|-------|---|---|---|---|---|
| row 0 | 0 | 1 | 1 | 1 | 0 |
| row 1 | 1 | 0 | 1 | 1 | 0 |
| row 2 | 1 | 1 | 0 | 1 | 1 |
| row 3 | 1 | 1 | 1 | 0 | 1 |
| row 4 | 0 | 0 | 1 | 1 | 0 |

Note: The number of values of 1 in the 2D table is twice the number of edges in the connectivity graph

Note: `char` is also an integer type in C, one byte per value

Read a connectivity graph from file (version 1)

```
void read_graph_from_file1 (char *filename, int *N, char ***table2D)
{
    int num_nodes, num_edges, i, w1, w2; char line[100], str1[50], str2[50];
    char **table;

    FILE *fp = fopen(filename, "r");
    fgets (line, 100, fp); // skip line 1
    fgets (line, 100, fp); // skip line 2
    fgets (line, 100, fp);
    sscanf (line, "# %s %d %s %d", str1, &num_nodes, str2, &num_edges);
    fgets (line, 100, fp); // skip line 4

    table = (char**)malloc(num_nodes*sizeof(char*));
    table[0] = (char*)calloc(num_nodes*num_nodes, sizeof(char));
    for (i=1; i<num_nodes; i++)
        table[i] = &(table[0][i*num_nodes]);

    for (i=0; i<num_edges; i++) {
        fgets (line, 100, fp); sscanf (line, "%d %d", &w1, &w2);
        if (w1!=w2 && w1>=0 && w1<num_nodes && w2>=0 && w2<num_nodes)
            table[w2][w1] = table[w1][w2] = 1;
    }
    fclose(fp);

    *N = num_nodes; *table2D = table;
}
```

Compressed row storage (CRS) format for connectivity graph

Implemented as two arrays of integer values, for example,

row_ptr: 0, 3, 6, 10, 14, 16

col_idx: 1, 2, 3, 0, 2, 3, 0, 1, 3, 4, 0, 1, 2, 4, 2, 3

Length of row_ptr: number of nodes in the connectivity graph + 1


Length of col_idx: twice the number of edges

Read a connectivity graph from file (version 2)

- Get the info about the number of nodes
- Allocate array `row_ptr`
- Read through the input file **in two rounds**
- Round 1:
 - Count how many nearest neighbors each node has
 - Prepare the values inside array `row_ptr`
- Allocate array `col_idx`
- Round 2:
 - Parse the info of each edge, insert *two entries* suitably into `col_idx`

SNN graph

For each pair of nearest neighbors u and v , the number of their *shared nearest neighbors* (SNNs) is defined as the number of other nodes that are directly connected with both u and v .



Naive implementation of create_SNN_graph1

```
void create_SNN_graph1 (int N, char **table2D, int ***SNN_table)
{
    int i,j,k,snn;

    int **s_table = (int**)malloc(N*sizeof(int*));
    s_table[0] = (int*)malloc(N*N*sizeof(int));
    for (i=1; i<N; i++)
        s_table[i] = &(s_table[0][i*N]);

    for (i=0; i<N; i++)
        for (j=0; j<N; j++)
            if (table2D[i][j]==1) {
                snn = 0;
                for (k=0; k<N; k++)
                    if (table2D[i][k]==1 && table2D[j][k]==1)
                        ++snn;
                s_table[i][j] = snn;
            }
            else
                s_table[i][j] = 0;

    *SNN_table = s_table;
}
```

Two improvements of create_SNN_graph1

```
// ....  
  
for (i=0; i<N-1; i++) {  
    s_table[i][i] = 0;  
  
    for (j=i+1; j<N; j++)  
        if (table2D[i][j]==1) {  
            snn = 0;  
            for (k=0; k<N; k++)  
                snn += table2D[i][k]*table2D[j][k];  
            s_table[i][j] = s_table[j][i] = snn;  
        }  
    else  
        s_table[i][j] = s_table[j][i] = 0;  
}  
s_table[N-1][N-1] = 0;
```

- Use symmetry of SNN pairs (saving 50% computation)
- Use the integer values (0 & 1) in table2D directly in computation

Naive implementation of create_SNN_graph2

```
void create_SNN_graph2 (int N, int *row_ptr, int *col_idx, int **SNN_val)
{
    int u,v,j, len1,len2;
    int *s_val = (int*)malloc(row_ptr[N]*sizeof(int));

    for (u=0; u<N; u++) {
        len1 = row_ptr[u+1]-row_ptr[u];
        for (j=row_ptr[u]; j<row_ptr[u+1]; j++) {
            v = col_idx[j]; len2 = row_ptr[v+1]-row_ptr[v];
            s_val[j] = num_shared_indices(len1,&(col_idx[row_ptr[u]]),
                                         len2,&(col_idx[row_ptr[v]]));
        }
    }
    *SNN_val = s_val;
}

int num_shared_indices (int len1, int *indices1, int len2, int *indices2) {
    int k,m,idx,num=0;
    for (k=0; k<len1; k++) {
        idx = indices1[k];
        for (m=0; m<len2; m++)
            if (indices2[m]==idx)
                ++num;
    }
    return num;
}
```

Improvements of create_SNN_graph2

- Sort the `col_idx` segment for each row beforehand, using for example *quick-sort* (“extra work” , but it’s worth it!)
- Counting shared indices between two sorted index lists can make use of *binary search* (which is much faster)
- Use the symmetry of SNN pairs (saving 50% computation)

```
for (u=0; u<N; u++) {
    len1 = row_ptr[u+1]-row_ptr[u];
    for (j=row_ptr[u]; j<row_ptr[u+1]; j++) {
        v = col_idx[j];
        if (v>u) { // using the symmetry of SNN pairs
            len2 = row_ptr[v+1]-row_ptr[v];
            s_val[j] = num_shared_indices_sorted(len1,&(col_idx[row_ptr[u]]),
                                                len2,&(col_idx[row_ptr[v]]));
            // find the position of u inside the col_idx segment of row v
            // u_pos = binary_search ....
            s_val[u_pos] = s_val[j];
        }
    }
}
```

OpenMP parallelization

```
#pragma omp parallel for private(j,k,snn)
for (i=0; i<N-1; i++) {
    s_table[i][i] = 0;

    for (j=i+1; j<N; j++)
        if (table2D[i][j]==1) {
            snn = 0;
            for (k=0; k<N; k++)
                snn += table2D[i][k]*table2D[j][k];
            s_table[i][j] = s_table[j][i] = snn;
        }
    else
        s_table[i][j] = s_table[j][i] = 0;
}
```

- Straightforward OpenMP parallelization of the outer-most loop (no danger for race conditions)
- Important to mark some of the variables as *thread-private*
- Load imbalance is a potential problem, should use small *chunksize* or *dynamic* scheduler
- Same parallelization strategy for `create_SNN_graph2`

“Clustering” (only relevant for IN4200 students)

```
void check_node (int node_id, int tau, int N,  
                int *row_ptr, int *col_idx, int *SNN_val)  
{  
    char *in_cluster = (char*)calloc(N, sizeof(char));  
    in_cluster[node_id] = 1; // mark the node  
  
    traverse_one_row (tau, node_id, row_ptr, col_idx, SNN_val, in_cluster);  
  
    // print nodes whose "in_cluster" values are non-zero  
    // ....  
}  
  
// a recursive function for depth-first search  
void traverse_one_row (int tau, int row_nr, int *row_ptr, int *col_idx,  
                      int *SNN_val, char *in_cluster)  
{  
    for (int i=row_ptr[row_nr]; i<row_ptr[row_nr+1]; i++)  
        if (in_cluster[col_idx[i]]==0 && SNN_val[i]>=tau) {  
            in_cluster[col_idx[i]] = 1; // mark the node  
            traverse_one_row (tau, col_idx[i], row_ptr, col_idx,  
                             SNN_val, in_cluster);  
        }  
}
```